

Oracle - Advanced SQL

Joins, Aggregations, Subqueries & Set Operators

By

Narasimha Rao T

Microsoft.Net FSD Trainer

Professional Development Trainer

tnrao.trainer@gmail.com

1. What are Joins?

Definition:

Joins are used to combine rows from two or more tables based on a related column between them. They allow you to retrieve data from multiple tables in a single query.

Why We Need Joins:


- Data is often spread across multiple tables (normalization)
- Avoids data redundancy
- Maintains data integrity
- Enables complex data retrieval

2. Advantages of Using Joins


Benefits:

1. **Data Integration:** Combine related data from multiple tables
2. **Performance:** Single query vs multiple queries
3. **Data Consistency:** Maintains relationships between tables
4. **Flexibility:** Complex data retrieval capabilities
5. **Efficiency:** Reduces application-level processing


Self Join




Right Join




Full Join




Inner Join



Cross Join



Left Join



3. Types of JOINS

INNER JOIN

Returns only the rows that have matching values in both tables.

```
-- Basic INNER JOIN
SELECT e.employee_id, e.first_name, e.last_name, d.department_name
FROM employees e
INNER JOIN departments d ON e.department_id = d.department_id;

-- Equivalent using WHERE clause (older syntax)
SELECT e.employee_id, e.first_name, e.last_name, d.department_name
FROM employees e, departments d
WHERE e.department_id = d.department_id;
```

LEFT OUTER JOIN

Returns all rows from the left table and matched rows from the right table. Unmatched rows from right table contain NULL.

```
-- LEFT JOIN - All employees, even if no department
SELECT e.employee_id, e.first_name, e.last_name, d.department_name
FROM employees e
LEFT JOIN departments d ON e.department_id = d.department_id;

-- LEFT JOIN with WHERE to find employees without departments
SELECT e.employee_id, e.first_name, e.last_name
FROM employees e
LEFT JOIN departments d ON e.department_id = d.department_id
WHERE d.department_id IS NULL;
```

RIGHT OUTER JOIN

Returns all rows from the right table and matched rows from the left table. Unmatched rows from left table contain NULL.

```
-- RIGHT JOIN - All departments, even if no employees
SELECT d.department_name, e.employee_id, e.first_name, e.last_name
FROM employees e
RIGHT JOIN departments d ON e.department_id = d.department_id;

-- RIGHT JOIN with WHERE to find departments without employees
SELECT d.department_name
FROM employees e
RIGHT JOIN departments d ON e.department_id = d.department_id
WHERE e.employee_id IS NULL;
```

FULL OUTER JOIN

Returns all rows when there's a match in either left or right table.

```
-- FULL OUTER JOIN
SELECT e.employee_id, e.first_name, d.department_name
FROM employees e
FULL OUTER JOIN departments d ON e.department_id = d.department_id;
```


SELF JOIN

Joins a table to itself. Useful for hierarchical data or comparing rows within the same table.

```
-- Find employees and their managers
SELECT e.employee_id, e.first_name AS employee_name,
       m.first_name AS manager_name
FROM employees e
LEFT JOIN employees m ON e.manager_id = m.employee_id;

-- Find employees in the same department
SELECT e1.first_name AS employee1, e2.first_name AS employee2, e1.department_id
FROM employees e1
JOIN employees e2 ON e1.department_id = e2.department_id
WHERE e1.employee_id < e2.employee_id; -- Avoids duplicates and self-comparison
```

CROSS JOIN

Returns Cartesian product of both tables (every row from first table combined with every row from second table).

```
-- CROSS JOIN
SELECT e.first_name, d.department_name
FROM employees e
CROSS JOIN departments d;
```

4. What are Aggregations?

Definition:

Aggregation operations process multiple rows to return a single summary value. They are used to perform calculations on sets of rows.

Common Use Cases:

- Counting records
- Calculating averages
- Finding maximum/minimum values
- Summing values
- Grouping data for analysis

5. Aggregate Functions and GROUP BY

Common Aggregate Functions:

```
-- COUNT - Count number of rows
SELECT COUNT(*) FROM employees;
SELECT COUNT(department_id) FROM employees; -- Excludes NULLs
SELECT COUNT(DISTINCT department_id) FROM employees;

-- SUM - Calculate total
SELECT SUM(salary) FROM employees;
SELECT SUM(salary) FROM employees WHERE department_id = 50;

-- AVG - Calculate average
SELECT AVG(salary) FROM employees;
SELECT AVG(NVL(salary, 0)) FROM employees; -- Handle NULLs
```

```
-- MAX/MIN - Find maximum/minimum values  
SELECT MAX(salary), MIN(salary) FROM employees;  
SELECT MAX(hire_date), MIN(hire_date) FROM employees;
```

GROUP BY Clause

Groups rows that have the same values into summary rows.

```
-- Basic GROUP BY
SELECT department_id, COUNT(*) as employee_count
FROM employees
GROUP BY department_id;

-- Multiple columns GROUP BY
SELECT department_id, job_id, COUNT(*) as employee_count
FROM employees
GROUP BY department_id, job_id;
```

```
-- GROUP BY with aggregate functions
SELECT department_id,
       COUNT(*) as employee_count,
       AVG(salary) as avg_salary,
       MAX(salary) as max_salary,
       MIN(salary) as min_salary
FROM employees
GROUP BY department_id;
```

HAVING Clause

Filters groups based on aggregate conditions (WHERE filters rows, HAVING filters groups).

```
-- HAVING with aggregate condition
SELECT department_id, COUNT(*) as employee_count
FROM employees
GROUP BY department_id
HAVING COUNT(*) > 5;
```

```
-- Multiple HAVING conditions
SELECT department_id, AVG(salary) as avg_salary
FROM employees
GROUP BY department_id
HAVING AVG(salary) > 5000 AND COUNT(*) > 3;
```



```
-- WHERE and HAVING together
SELECT department_id, AVG(salary) as avg_salary
FROM employees
WHERE hire_date > DATE '2020-01-01'
GROUP BY department_id
HAVING AVG(salary) > 6000;
```

6. Working with Filtering and Sorting

Advanced Filtering Techniques:

```
-- Multiple conditions
SELECT * FROM employees
WHERE salary BETWEEN 5000 AND 10000
      AND department_id IN (10, 20, 30)
      AND hire_date > DATE '2015-01-01';

-- Pattern matching with LIKE
SELECT * FROM employees
WHERE first_name LIKE 'A%';           -- Starts with A
WHERE first_name LIKE '%a%';         -- Contains 'a'
WHERE first_name LIKE '_a%';         -- Second letter is 'a'
WHERE first_name LIKE 'J_n%';        -- Starts with J, third letter n
```

```
-- NULL handling
SELECT * FROM employees
WHERE commission_pct IS NOT NULL;

-- EXISTS with subquery
SELECT * FROM departments d
WHERE EXISTS (
    SELECT 1 FROM employees e
    WHERE e.department_id = d.department_id
);
```

Advanced Sorting:

```
-- Basic sorting
SELECT first_name, last_name, salary
FROM employees
ORDER BY salary DESC;

-- Multiple column sorting
SELECT department_id, first_name, last_name, salary
FROM employees
ORDER BY department_id ASC, salary DESC;

-- Sorting by expression
SELECT first_name, last_name, salary, commission_pct
FROM employees
ORDER BY NVL(commission_pct, 0) DESC;
```

```
-- Sorting by column position
SELECT first_name, last_name, salary
FROM employees
ORDER BY 3 DESC, 2 ASC; -- Sort by 3rd column (salary), then 2nd (last_name)

-- CASE statement in ORDER BY
SELECT first_name, last_name, department_id, salary
FROM employees
ORDER BY
    CASE
        WHEN department_id = 50 THEN 1
        WHEN department_id = 60 THEN 2
        ELSE 3
    END,
    salary DESC;
```

7. What are Subqueries?

Definition:

A subquery (inner query or nested query) is a query within another SQL query. It's used to return data that will be used in the main query.

Types of Subqueries:

1. **Single-row subquery:** Returns one row
2. **Multiple-row subquery:** Returns multiple rows

8. Examples of Subqueries

Single-Row Subqueries:

```
-- Employees with salary greater than average
SELECT first_name, last_name, salary
FROM employees
WHERE salary > (SELECT AVG(salary) FROM employees);

-- Employee with highest salary
SELECT first_name, last_name, salary
FROM employees
WHERE salary = (SELECT MAX(salary) FROM employees);
```

Multiple-Row Subqueries:

```
-- Using IN
SELECT first_name, last_name, department_id
FROM employees
WHERE department_id IN (
    SELECT department_id
    FROM departments
    WHERE location_id = 1700
);
```

```
-- Using ANY/SOME
SELECT first_name, last_name, salary
FROM employees
WHERE salary > ANY (
    SELECT salary
    FROM employees
    WHERE department_id = 50
);
```


Correlated Subqueries:

```
-- Employees whose salary is greater than department average
SELECT e1.first_name, e1.last_name, e1.salary, e1.department_id
FROM employees e1
WHERE salary > (
    SELECT AVG(salary)
    FROM employees e2
    WHERE e2.department_id = e1.department_id
);
```

```
-- Departments that have employees
SELECT department_name
FROM departments d
WHERE EXISTS (
    SELECT 1
    FROM employees e
    WHERE e.department_id = d.department_id
);
```

9. What are Set Operators?

Definition:

Set operators combine the results of two or more SELECT statements. All SELECT statements must have the same number of columns and compatible data types.

Types of Set Operators:

UNION

Combines results and removes duplicates.

```
-- Employees and managers from different tables
SELECT employee_id, first_name, last_name, 'Employee' as type
FROM employees
UNION
SELECT manager_id, first_name, last_name, 'Manager' as type
FROM managers;
```

UNION ALL

Combines results and keeps all duplicates.

```
-- All locations including duplicates  
SELECT city FROM old_locations  
UNION ALL  
SELECT city FROM new_locations;
```

INTERSECT

Returns only common rows from both queries.

```
-- Products available in both warehouses
SELECT product_id FROM warehouse_a
INTERSECT
SELECT product_id FROM warehouse_b;
```

MINUS (EXCEPT in other databases)

Returns rows from first query that are not in second query.

```
-- Products only in warehouse A
SELECT product_id FROM warehouse_a
MINUS
SELECT product_id FROM warehouse_b;
```

Q & A

Narasimha Rao T

tnrao.trainer@gmail.com